



3

MPI Parallel Programming

Lecture 1

Dr Abid Mehmood
[\(abid.mehmood@dsu.edu\)](mailto:(abid.mehmood@dsu.edu))

Dakota State University

Learning Objectives

- Understanding how MPI programs execute
- Familiarity with fundamental MPI functions




Outline


- Message-passing model
- Message Passing Interface (MPI)
- Coding MPI programs
- Compiling and running MPI programs
- Benchmarking MPI programs



Message-passing Model



Task/Channel vs. Message-passing



- Task/channel
 - Task
 - Explicit channels
- Message-passing
 - Process
 - Any-to-any communication



Process Characteristics in Message-Passing Model

- Number of processes is **specified at launch time**
- **Remains constant** throughout execution of program
- All processes execute the **same program**
- Each process has **unique ID** number
- Processes **alternate** between performing computations and communicating with each other.



Advantages of Message-Passing Model

- Gives programmer ability to manage the memory hierarchy
 - Maximize local computations while minimizing communications
- Portability to many architectures
- Easier to create a deterministic program
- Simplifies debugging



The Message Passing Interface

- Late 1980s: vendors had unique libraries
- 1989: Parallel Virtual Machine (PVM) developed at Oak Ridge National Lab
- 1992: Work on MPI standard began
- 1994: Version 1.0
- 1997: Version 2.0
- 2012: Version 3.0
- 2021: Version 4.0
- 2023: Version 4.1
- <https://www.mpi-forum.org/>
- The Open MPI Project is an open source MPI implementation developed and maintained by a consortium of academic, research, and industry partners.
- <https://www.open-mpi.org/>



```
1 #include <stdio.h>
2 #include <mpi.h>
3
4 int main(int argc, char *argv[]) {
5     int numprocs, rank, namelen;
6     char processor_name[MPI_MAX_PROCESSOR_NAME];
7
8     MPI_Init(&argc, &argv);
9     MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
10    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
11    MPI_Get_processor_name(processor_name, &namelen);
12
13    printf("Process %d on %s out of %d\n", rank, processor_name, numprocs);
14
15    MPI_Finalize();
16 }
```

```
mpiuser@Washington:~/nfs/CSC718/openMPI/hello$ mpicc hello.c -o hello_mpi
mpiuser@Washington:~/nfs/CSC718/openMPI/hello$ mpirun -np 4 -machinefile machine
file.dsu ./hello_mpi
Process 0 on Lincoln out of 4
Process 2 on Roosevelt out of 4
Process 1 on Washington out of 4
Process 3 on Jefferson out of 4
mpiuser@Washington:~/nfs/CSC718/openMPI/hello$
```




Circuit Satisfiability Problem (CSAT) for MPI

- Demonstrates key MPI principles like process communication, parallel task distribution, and performance scaling.
- A well-known NP-complete problem
 - No known algorithms to solve in polynomial time
- You determine if there is a combination of inputs that causes the output of a Boolean circuit to be true ('1').
- Requires testing numerous input combinations.
- Exhaustive search for solutions is highly parallelizable:
 - Each possible input combination can be tested independently.
 - The problem fits naturally with MPI's message-passing paradigm.



Circuit Satisfiability




Not satisfied

To determine if some combination of inputs cause the output of the circuit to be 1

16 inputs \Rightarrow 65,536 combinations to test



Partitioning: Functional Decomposition



Embarrassingly parallel: No channels between tasks

Embarrassingly Parallel:

Tasks that can be easily separated into independent sub-tasks.

Minimal communication between tasks.

Ideal for distributed systems and parallel programming.

Examples: Image processing, data encryption, large-scale simulations.



Agglomeration and Mapping in Parallel Algorithms

Agglomeration

- Grouping smaller tasks into larger units to minimize communication overhead and balance the workload.
- Properties of parallel algorithm
 - Fixed number of tasks (determined upfront)
 - No communications between tasks (no data exchange)
 - Time needed per task is variable (some may take longer)

Task mapping strategy

- Decides how to distribute (map) tasks to processors
- Use cyclic mapping to assign tasks to processors in a round-robin fashion to balance the workload.



Cyclic (Interleaved) Allocation

- Assume p processes and N tasks
- Each process gets every p -th piece of work
 - process i gets tasks numbered: $i, i + p, i + 2p, i + 3p, \dots$ until all tasks are assigned
- Example: 5 processes and 12 pieces of work
 - P0: 0, 5, 10
 - P1: 1, 6, 11
 - P2: 2, 7
 - P3: 3, 8
 - P4: 4, 9

When the number of tasks is not perfectly divisible by the number of processes, some processes will end up with more tasks than others.



Exercise

- Assume n pieces of work, p processes, and cyclic allocation
- What is the most pieces of work any process has?
- What is the least pieces of work any process has?
- How many processes have the most pieces of work?



CSAT: Summary of Program Design

- Program will consider all 65,536 combinations of 16 boolean inputs
- Combinations allocated in cyclic fashion to processes
- Each process examines each of its combinations
- If it finds a satisfiable combination, it will output it

The example encapsulates key MPI ideas:

- Initialization (MPI_Init)*
- Rank-based task division*
- Independent local computation*
- Optional result communication*
- Finalization (MPI_Finalize)*



Key MPI Terminology Summary

- **Process Rank:** Unique identifier for each process.
- **Communicator:** Group of processes involved in message passing.
- **Cyclic Allocation:** Distributes work evenly across processes.
- **Collective Communication:** Communication involving multiple processes.
- **Synchronization:** Ensuring processes are in sync.





3

MPI Parallel Programming

Lecture 2

Dr Abid Mehmood
[\(abid.mehmood@dsu.edu\)](mailto:(abid.mehmood@dsu.edu))

Dakota State University

Objectives

- Parallelizing a task across multiple processes
- Enabling communication between processes
- Timing and performance in parallel programs



circuit1.c

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main (int argc, char *argv[])
5 {
6     int i;
7     int id;
8     int p;
9     void check_circuit (int, int);
10    MPI_Init (&argc, &argv);
11    MPI_Comm_rank (MPI_COMM_WORLD, &id);
12    MPI_Comm_size (MPI_COMM_WORLD, &p);
13    for (i = id; i < 65536; i += p)
14        check_circuit (id, i);
15
16    printf ("Process %d is done\n", id);
17    fflush (stdout);
18    MPI_Finalize();
19    return 0;
20 }
```





```
22 /* Return 1 if 'i'th bit of 'n' is 1; 0 otherwise */
23 #define EXTRACT_BIT(n,i) ((n&(1<<i))?1:0)
24
25 void check_circuit (int id, int z) {
26     int v[16];          /* Each element is a bit of z */
27     int i;
28
29     for (i = 0; i < 16; i++) v[i] = EXTRACT_BIT(z,i);
30
31     if ((v[0] || v[1]) && (!v[1] || !v[3]) && (v[2] || v[3])
32         && (!v[3] || !v[4]) && (v[4] || !v[5])
33         && (v[5] || !v[6]) && (v[5] || v[6])
34         && (v[6] || !v[15]) && (v[7] || !v[8])
35         && (!v[7] || !v[13]) && (v[8] || v[9])
36         && (v[8] || !v[9]) && (!v[9] || !v[10])
37         && (v[9] || v[11]) && (v[10] || v[11])
38         && (v[12] || v[13]) && (v[13] || !v[14])
39         && (v[14] || v[15])) {
40         //printf("satisfied for number: %d\n", z);
41         printf ("%d) %d%d%d%d%d%d%d%d%d%d%d%d%d%d%d\n", id,
42                 v[0],v[1],v[2],v[3],v[4],v[5],v[6],v[7],v[8],v[9],
43                 v[10],v[11],v[12],v[13],v[14],v[15]);
44         fflush (stdout);
45     }
46 }
```

circuit1.c

Local Variables

```
int main (int argc, char *argv[]) {  
    int i;  
    int id; /* Process rank */  
    int p; /* Number of processes */  
    void check_circuit (int, int);
```

- Include `argc` and `argv`: they are needed to initialize MPI
- One copy of every variable for each process running this program



Initialize MPI

```
MPI_Init (&argc, &argv) ;
```

- First MPI function called by each process
- Not necessarily first executable statement
- Allows system to do any necessary setup



Types of Communication in Parallel Programming

- **Point-to-point communication:** Direct communication between two processes.
 - MPI functions: MPI_Send, MPI_Recv.
- **Collective communication:** Involves groups of processes.
 - Examples: Broadcast, reduce, scatter, and gather.
 - MPI functions: MPI_Bcast, MPI_Reduce, MPI_Scatter, MPI_Gather.
- **Synchronization:** Ensuring processes are aligned before continuing.
 - MPI function: MPI_Barrier.




Communicator and MPI_COMM_WORLD

- **Communicator:** Object that provides message-passing environment for processes
- Every communicator in MPI has a name.
- **MPI_COMM_WORLD**
 - Default communicator
 - Includes all processes
- Each process knows its rank **within a communicator**
 - Could have **different ranks** in different communicators
- Possible to create new communicators
 - Will cover later in course



Communicator and MPI_COMM_WORLD



Determine Number of Processes

```
MPI_Comm_size (MPI_COMM_WORLD, &p) ;
```

- First argument is communicator
- Number of processes returned through second argument




Determine Process Rank

```
MPI_Comm_rank (MPI_COMM_WORLD, &id) ;
```

- First argument is communicator
- Process rank (in range 0, 1, ..., p-1) returned through second argument



Replication of Automatic Variables



Each process has its own copy of variables declared within a function (automatic variables).

Automatic variables are not shared among processes.

Recall that MPI programs are for distributed memory systems, not shared memory.



What about External Variables?

```
int total;
```

Global variable

```
int main (int argc, char *argv[]) {  
    int i;  
    int id;  
    int p;  
    ...
```

- Where is variable **total** stored?



What about External Variables?

```
int total;           Global variable
```

```
int main (int argc, char  
*argv[]) {
```

...

- Global variables are accessible to all functions [within the same source file](#) (or other files if declared as extern)
- [Each process](#) has its own copy of global variables ([not shared](#)).
- To share, MPI communication function will be used (e.g. MPI_Send, MPI_Bcast)



Cyclic Allocation of Work

```
for (i = id; i < 65536; i += p)
    check_circuit (id, i);
```

- The loop determines **which combinations** each process will check.
- Parallelism is outside the function **check_circuit**
- It can be an ordinary, sequential function.
- The MPI logic only handles **which data** each process **should process**, not **how the computation itself is done**.



Shutting Down MPI

```
MPI_Finalize();
```

- Call after all other MPI library calls
- Allows system to free up MPI resources



Compiling MPI Programs

```
mpicc -O -o foo foo.c
```

- **mpicc**: script to compile and link C+MPI programs
- **Flags**: same meaning as C compiler
 - **-O** — optimize
 - **-o <file>** — where to put executable



Running MPI Programs

- **mpirun -np <p> <exec> <arg1> ...**
 - **-np <p>** — number of processes
 - **<exec>** — executable
 - **<arg1> ...** — command-line arguments



Specifying Host Processors

- `mpirun -np 2 -machinefile machinefile.dsuhello`
- (default) File `.mpi-machines` in home directory lists host processors in order of their use
- Example machinefile contents
 - band01.cs.ppu.edu
 - band02.cs.ppu.edu
 - band03.cs.ppu.edu
 - band04.cs.ppu.edu

If `-np` exceeds the number of listed machines, MPI cycles through the list again (depending on configuration).



Enabling Remote Logins

- MPI needs to be able to initiate processes on other processors without supplying a password
- Each processor in group must list all other processors in its **.rhosts file** [hostname username] e.g.,
band01.cs.ppu.edu student
 - band02.cs.ppu.edu student
 - band03.cs.ppu.edu student
 - band04.cs.ppu.edu student



Deciphering Output

- Output order only **partially reflects** order of output events inside parallel computer
- If process A prints two messages, first message will appear before second (**guaranteed intra-process ordering**)
- If process A calls **printf** before process B, **there is no guarantee** process A's message will appear before process B's message

Impl Tip: If ordered output is required for readability (e.g., by process rank):

- Use **MPI_Barrier()** before printing, or
- Have only the root process print results after gathering them.



```
(base) abid@MacBookAir circuit % mpirun -np 1 circuit1
0) 1010111110011001
0) 0110111110011001
0) 1110111110011001
0) 1010111111011001
0) 0110111111011001
0) 1110111111011001
0) 10101111110111001
0) 01101111110111001
0) 11101111110111001
Process 0 is done
```

```
(base) abid@MacBookAir circuit % mpirun -np 2 circuit1
0) 0110111110011001
0) 0110111111011001
0) 01101111110111001
1) 1010111110011001
1) 1110111110011001
1) 1010111111011001
1) 1110111111011001
1) 10101111110111001
1) 11101111110111001
Process 0 is done
Process 1 is done
```



```
(base) abid@MacBookAir circuit % mpirun -np 3 circuit1
0) 0110111110011001
0) 1110111111011001
0) 1010111110111001
2) 1010111110011001
2) 0110111111011001
2) 11101111110111001
1) 11101111110011001
1) 1010111111011001
1) 01101111110111001
Process 0 is done
Process 1 is done
Process 2 is done
```

```
(base) abid@MacBookAir circuit % mpirun -np 4 circuit1
3) 1110111110011001
3) 1110111111011001
3) 11101111110111001
1) 1010111110011001
1) 1010111111011001
1) 10101111110111001
2) 0110111110011001
2) 0110111111011001
2) 01101111110111001
Process 0 is done
Process 3 is done
Process 2 is done
Process 1 is done
```



Enhancing the Program

- We want to find total number of solutions
- Incorporate sum-reduction into program
- Reduction is a **collective communication**



Modifications

- Modify function **check_circuit**
 - Return 1 if circuit satisfiable with input combination
 - Return 0 otherwise
- Each process keeps local count of satisfiable circuits it has found
- Perform reduction after **for** loop



New Declarations and Code

```
int count; /* Local sum */  
int global_count; /* Global sum */  
int check_circuit (int, int);  
  
count = 0;  
for (i = id; i < 65536; i += p)  
    count += check_circuit (id, i);
```



Definition of MPI_Reduce()

```
int MPI_Reduce(  
    void *operand, /* address of send buffer */  
    void *result, /* address of receive buffer (only at root) */  
    int count,   /* number of elements */  
    MPI_Datatype type, /* data type of elements */  
    MPI_Op operator, /* reduction operation */  
    int root,     /* rank of process receiving result */  
    MPI_Comm comm /* communicator */  
);
```

```
double local_sum, global_sum;  
MPI_Reduce(&local_sum, &global_sum, 1, MPI_DOUBLE,  
MPI_SUM, 0, MPI_COMM_WORLD);
```

All processes send their local double values, and process 0 receives the combined sum.



MPI_Datatype Options

- MPI_CHAR
- MPI_DOUBLE
- MPI_FLOAT
- MPI_INT
- MPI_LONG
- MPI_LONG_DOUBLE
- MPI_SHORT
- MPI_UNSIGNED_CHAR
- MPI_UNSIGNED
- MPI_UNSIGNED_LONG
- MPI_UNSIGNED_SHORT



MPI_Op Options

- **MPI_MAX** maximum
- **MPI_MIN** minimum
- **MPI_SUM** sum
- **MPI_PROD** product
- **MPI LAND** logical and
- **MPI_BAND** bit-wise and
- **MPI_LOR** logical or
- **MPI_BOR** bit-wise or
- **MPI_LXOR** logical xor
- **MPI_BXOR** bit-wise xor
- **MPI_MAXLOC** max value and location (rank)
- **MPI_MINLOC** min value and location



Our Call to MPI_Reduce()

```
MPI_Reduce (&count,  
            &global_count,  
            1,  
            MPI_INT,  
            MPI_SUM,  
            0,  
            MPI_COMM_WORLD) ;
```

Process 0 gets
the result

```
if (!id) printf ("There are %d different solutions\n",  
    global_count);
```



```
4 int main(int argc, char *argv[])
5 {
6     int i;
7     int id;
8     int p;
9
10    int count; /* Local sum */
11    int global_count; /* Global sum */
12    int check_circuit (int, int);
13
14    MPI_Init (&argc, &argv);
15    MPI_Comm_rank (MPI_COMM_WORLD, &id);
16    MPI_Comm_size (MPI_COMM_WORLD, &p);
17
18    count = 0;
19    for (i = id; i < 65536; i += p)
20        count += check_circuit (id, i);
```



```
21  /* aggregate data from all processes in a communicator and
22   reduce it to a single result*/
23   MPI_Reduce (&count, //send buffer (local data)
24             &global_count, // receive buffer (global data)
25             1, //how many elements per process are to be reduced (1: count)
26             MPI_INT, // type of element to be reduced
27             MPI_SUM, // op to be performed on data
28             0, // in which process data will be stored
29             MPI_COMM_WORLD);
30
31   printf ("Process %d is done\n", id);
32   fflush (stdout);
33   MPI_Finalize();
34
35   if (!id) printf ("There are %d different solutions\n",
36                   global_count);
37
38   return 0;
39 }
40
41 /* Return 1 if 'i'th bit of 'n' is 1; 0 otherwise */
42 #define EXTRACT_BIT(n,i) ((n&(1<<i))?1:0)
--
```



```
44 int check_circuit (int id, int z) {
45     int v[16];          /* Each element is a bit of z */
46     int i;
47
48     for (i = 0; i < 16; i++) v[i] = EXTRACT_BIT(z,i);
49
50     if ((v[0] || v[1]) && (!v[1] || !v[3]) && (v[2] || v[3])
51         && (!v[3] || !v[4]) && (v[4] || !v[5])
52         && (v[5] || !v[6]) && (v[5] || v[6])
53         && (v[6] || !v[15]) && (v[7] || !v[8])
54         && (!v[7] || !v[13]) && (v[8] || v[9])
55         && (v[8] || !v[9]) && (!v[9] || !v[10])
56         && (v[9] || v[11]) && (v[10] || v[11])
57         && (v[12] || v[13]) && (v[13] || !v[14])
58         && (v[14] || v[15])) {
59         printf ("%d %d%d%d%d%d%d%d%d%d%d%d%d\n", id,
60                 v[0],v[1],v[2],v[3],v[4],v[5],v[6],v[7],v[8],v[9],
61                 v[10],v[11],v[12],v[13],v[14],v[15]);
62         fflush (stdout);
63         return 1;
64     }
65
66     return 0;
67 }
```



Other MPI Collective Communication Functions

- **MPI_Gather:** Collects values from all processes to a root process.
- **MPI_Scatter:** Distributes chunks of data from the root process to all others.
- **MPI_Bcast:** Broadcasts data from the root process to all other processes.



```
(base) abid@MacBookAir circuit % mpirun -np 3 circuit2
2) 101011110011001
2) 011011111011001
2) 1110111110111001
0) 0110111110011001
0) 1110111111011001
0) 1010111110111001
1) 1110111110011001
1) 1010111111011001
1) 0110111110111001
Process 1 is done
Process 2 is done
Process 0 is done
There are 9 different solutions
```



Benchmarking the Program

- **MPI_Barrier** — process synchronization, blocks until all processes have reached this routine.
- **MPI_Wtick** — timer resolution, returns the resolution of [MPI_Wtime](#)
- **MPI_Wtime** — in seconds, returns an elapsed time on the calling processor.



Benchmarking Code

```
double elapsed_time;  
...  
MPI_Init (&argc, &argv);  
MPI_Barrier (MPI_COMM_WORLD);  
elapsed_time = - MPI_Wtime(); //start timer  
...  
MPI_Reduce (...);  
elapsed_time += MPI_Wtime(); //stop timer
```



```
4 int main(int argc, char *argv[])
5 {
6     int i;
7     int id;
8     int p;
9
10    double elapsed_time;
11    int count; /* Local sum */
12    int global_count; /* Global sum */
13    int check_circuit (int, int);
14
15    MPI_Init (&argc, &argv);
16    MPI_Barrier (MPI_COMM_WORLD);
17    elapsed_time = - MPI_Wtime();
18
19    MPI_Comm_rank (MPI_COMM_WORLD, &id);
20    MPI_Comm_size (MPI_COMM_WORLD, &p);
21
22    count = 0;
23    for (i = id; i < 65536; i += p)
24        count += check_circuit (id, i);
25
```



```
26     MPI_Reduce (&count,
27                 &global_count,
28                 1,
29                 MPI_INT,
30                 MPI_SUM,
31                 0,
32                 MPI_COMM_WORLD);
33     elapsed_time += MPI_Wtime();
34
35 // printf ("Process %d is done.\n", id);
36 // fflush (stdout);
37 MPI_Finalize();
38
39 if (!id)
40 {
41     printf ("Elapsed time: %f.\n", elapsed_time);
42     fflush (stdout);
43     printf ("There are %d different solutions\n", global_count);
44 }
45
46 return 0;
47 }
48
49 /* Return 1 if 'i'th bit of 'n' is 1; 0 otherwise */
50 #define EXTRACT_BIT(n,i) ((n&(1<<i))?1:0)
51
```



```
52 int check_circuit (int id, int z) {
53     int v[16];          /* Each element is a bit of z */
54     int i;
55
56     for (i = 0; i < 16; i++) v[i] = EXTRACT_BIT(z,i);
57
58     if ((v[0] || v[1]) && (!v[1] || !v[3]) && (v[2] || v[3])
59         && (!v[3] || !v[4]) && (v[4] || !v[5])
60         && (v[5] || !v[6]) && (v[5] || v[6])
61         && (v[6] || !v[15]) && (v[7] || !v[8])
62         && (!v[7] || !v[13]) && (v[8] || v[9])
63         && (v[8] || !v[9]) && (!v[9] || !v[10])
64         && (v[9] || v[11]) && (v[10] || v[11])
65         && (v[12] || v[13]) && (v[13] || !v[14])
66         && (v[14] || v[15])) {
67         // printf ("%d %d%d%d%d%d%d%d%d%d%d%d%d%d\n", id,
68         // v[0],v[1],v[2],v[3],v[4],v[5],v[6],v[7],v[8],v[9],
69         // v[10],v[11],v[12],v[13],v[14],v[15]);
70         // fflush (stdout);
71         return 1;
72     }
73
74     return 0;
75 }
```



Benchmarking Results

```
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ mpirun -np 4 -machinefile machinefile.dsu ./circuit3
Elapsed time: 0.001894.
There are 9 different solutions
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ mpirun -np 4 ./circuit3
Elapsed time: 0.027015.
There are 9 different solutions
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ █
```

0.001894 seconds (~1.894 milliseconds) on 4 processes across **multiple machines**.

0.027015 seconds (~27.015 milliseconds) on 4 processes on the **same machine**.



Interpreting Benchmarking Results

- Speedup: How much faster a parallel algorithm performs compared to a sequential one.

$\text{Speedup} = \text{Sequential Time} / \text{Parallel Time}$.

- Efficiency: How effectively parallelization is utilized.

$\text{Efficiency} = \text{Speedup} / \text{Number of Processors}$.

- Scalability: How well performance improves as more processors are added.

$\text{Scalability} = \text{Time with 1 processor} / \text{Time with n processors}$



Speedup Example

```
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ mpirun ./circuit3
Elapsed time: 0.003307.
There are 9 different solutions
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ mpirun -np 4 -machinefile machinefile.dsu ./circuit3
Elapsed time: 0.002584.
There are 9 different solutions
mpiuser@Roosevelt:~/nfs/CSC718/openMPI/circuit$ █
```

1 process: 0.003307 seconds (~3.307 ms)

4 processes: 0.002584 seconds (~2.584 ms)

Speedup = $0.003307 / 0.002584 = 1.279$

Parallel execution is ~1.279 times faster.

% Improvement = $(\text{Speedup} - 1) \times 100 = \sim 27.9\%$



Practical Applications of MPI

- **Scientific Simulations:** Weather modeling, climate prediction (e.g., NASA climate models).
- **Data Analytics:** Processing big data in machine learning systems.
- **Engineering Simulations:** Used in research for fluid dynamics, material science, etc.
- **High-Performance Computing:** MPI-based simulations run on systems like CERN or NASA for particle physics.

